

# Implementing the MapReduce architecture for small datasets

Dan Hughes

## Background

The MapReduce paradigm is a distributed data-processing framework for processing large amounts of data using a two-step execution of parallel functions. Designed at Google for processing log files [1], it enables commodity hardware to be added to a resource pool of machines.

The open-source implementation of MapReduce is a collection of software called Hadoop. This consists of a number of individual applications for resource management, file storage and other tasks. As a result, it has a significant set-up process and resource overheads.

This project presents an alternative platform to Hadoop, to be built entirely in JavaScript using the NodeJS runtime.

## Objectives

- The project should be a rapidly deployable platform for prototyping MapReduce algorithms on commodity hardware.
- Each individual node should automatically tune its own performance parameters, ensuring that slow hardware is able to contribute.
- Data should be processed with streaming operators, to ensure that memory limits on low-power devices are not exceeded.

## Algorithm

The general MapReduce algorithm follows this structure:

- Each node is assigned a subset of the data.
- Each node applies a map function in parallel to the data to build a list of key-value pairs.
- Each node is assigned all the values for a given key.
- Each node applies a reduce function in parallel to convert all the values for a given key to a single result.
- The results are output to file.


## Implementation

All nodes connect to the controller over a TCP WebSocket [2] to allow for real-time streaming of data. The controller generates a queue of work tasks to be processed by streaming the input file through a user-supplied stream transformation function. This allows the controller to always have work to be processed without overloading the machine's memory.

The worker nodes apply a map function to every value, store the value results in local memory, and send back the resulting key to the controller. Through this mechanism, the controller builds a list of which nodes have values for each key. The nodes also apply a user-supplied combination function, to perform local reduction of values. This allows for the network as a whole to act as a distributed object store, accessible to every node within it.

The controller then allocates each key to a node. Each node queries the distributed object store to find all the values for its assigned keys, applies the a reduction function, and returns these value to the controller.

## Scalability



## Applications

The purpose of the system is to allow users to prototype MapReduce algorithms on any available commodity hardware. It has been tested on systems ranging from a cheap Raspberry Pi cluster to powerful desktop systems. It has been applied to distributed classification problems such as K-Nearest Neighbour, and simpler algorithms such as calculating mutual friends for each pair of users in a large social graph.



### References

1. J. Dean and S. Ghemawat, 'MapReduce: Simplified data processing on large clusters', OSDI 2004
2. [RFC6445] I. Fette, "The WebSocket Protocol", RFC6445, December 2011

### Degree programme

BSc Computer Science with Industrial Year